

Isha Singh

Professor Irene Tsapara

MSDS 422 Practical Machine Learning

9 February 2025

### Module 05: Titanic

In 1912, the Titanic shipwreck resulted in the loss of about 1,502 lives out of the 2,224 passengers on board. This analysis used Kaggle's Titanic dataset to study survival rates based on factors such as age, gender, ticket price, and passenger class. The main goal within the research was to predict the number of survived through the help of machine learning models. The machine learning models that were focused on this fifth module were: Random Forest Regressor, Gradient Boosted Trees, and last but not least Extra Trees Regressor.

Before heading onto modeling and training, the first requirement was that the dataset be analyzed and prepared. Missing values were noticed in the dataset. The columns that had missing values were Age, Cabin, and Embarked. The missing values for Age were handled with the median value, as it was believed to be the most appropriate approach instead of replacing them with 0. Cabin was replaced with "unknown" wherever missing values were shown. Lastly, the column Embarked was replaced with the most common port (mode) based on the passenger's class.

A few feature engineering techniques were applied to achieve one mission—to improve the model's performance. A new column called Title was extracted from the Name. For example, in "Ms. Thomas," "Ms" was extracted using regex (regular expressions), which was learned in an NLP course taken previously. A column, Family Size, was created to understand the number of

members in the family to incorporate the understanding of whether the passenger was traveling alone or not. Age and Fare, which were numerical features, used Standard Scaling to maintain a consistent organization across the different variables.

Outliers in the dataset were further investigated with the help of a built-in method called the Interquartile Range (IQR), and the extreme values were removed so that they would not affect the models. Following that, the dataset was visualized to better understand its characteristics. A histogram was created to analyze the distribution of the Age column. It was noticeable that the distribution was right-skewed. Some observations from the visualization include: there were more young passengers compared to older ones, most passengers were adults around the median age, and no severe outliers were present. This suggests that age is a crucial factor and may help distinguish survival rates. Moreover, bar plots were created, revealing that passengers in first class had the highest chance of survival, followed by second class, with third class having the lowest survival rate. A box plot of Fare by passenger class was also generated, demonstrating that first-class tickets were the most expensive. Lastly, a heatmap was created to examine feature correlations, which revealed that Pclass and Fare were strongly related to the number of survivors.

Next, for evaluating the model, the dataset was divided into training and validation sets using the 5-fold cross-validation method, which was chosen to prevent overfitting. The models (Random Forest Regressor, Gradient Boosted Trees, and Extra Trees Regressor) were trained using Grid Search CV, which helped optimize the hyperparameters. The main hyperparameters used throughout the research were:

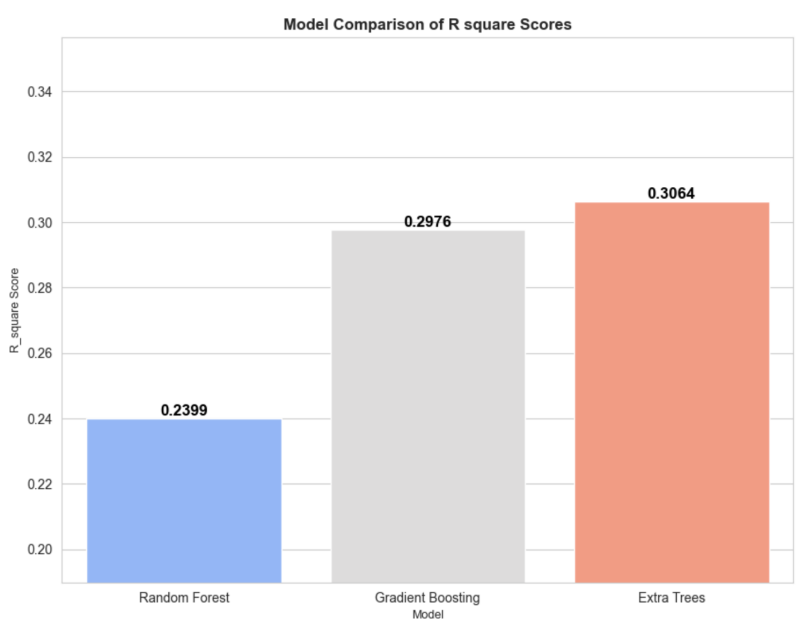

- number of trees within the model → `n_estimators`
- maximum depth of each tree → `max_depth`
- number of features that was considered per split → `max_features`
- minimum samples that are required in order to split a node → `min_samples_split`
- minimum samples in a leaf node → `min_samples_leaf`
- step size per iteration → `learning_rate`
- fraction of data used for boosting → `subsample`

For the assignment, one of the requirements was to use loss functions. For regression models, `squared_error` and `absolute_error` were used, as they are more suitable for regression tasks rather than classification methods. The models were also evaluated using Mean Squared Error (MSE) and R-Squared.

MODEL	MSE	R-SQUARED
<b>RANDOM FOREST</b>	0.152650	0.239911
<b>GRADIENT BOOSTING</b>	0.141070	0.297572
<b>EXTRA TREES</b>	0.139291	0.306430

Overall, according to the table created above, it is shown that the Extra Trees Regressor performed the best and most accurately, achieving the lowest MSE of 0.139291 and the highest R-squared value of 0.306430. The Gradient Boosting Regressor performed slightly worse than the Extra Trees Regressor but was comparatively better than the Random Forest Regressor. Unfortunately, the Random Forest Regressor performed the worst, with the highest MSE and the lowest R-squared value. To improve this, further hyperparameter tuning could be beneficial. Initially, I explored a wide range of hyperparameters, including more values for `max_depth` and

min\_samples\_split. However, I realized that this approach was taking over 30 minutes, so I decided to select a smaller range instead. Additionally, bar plots were also created to understand the MSE and R-squared which is similar to the table above but in graphical perspective.



Overall, from this research we can understand that the Extra Trees Regression was the best output and have performed the best primarily. For improvement, feature engineering or better hyperparameters is necessary.

## References

“Extratreesclassifier.” scikit. Accessed February 10, 2025.  
<https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>.

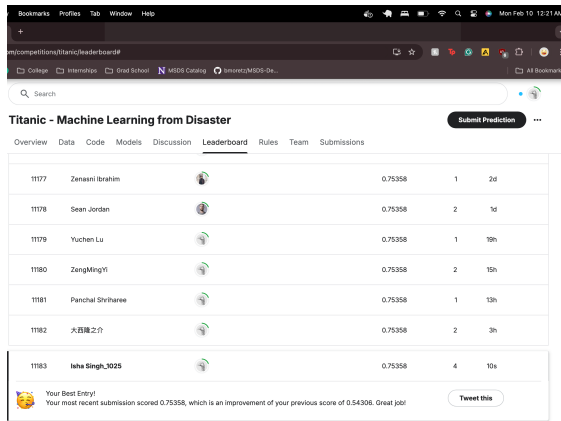
“Gradientboostingclassifier.” scikit. Accessed February 10, 2025. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>.

“Randomforestclassifier.” scikit. Accessed February 9, 2025. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>.

“Titanic - Machine Learning from Disaster.” Kaggle. Accessed February 9, 2025.  
<https://www.kaggle.com/competitions/titanic>.

## Kaggle Submissions

### Random Forest Regressor



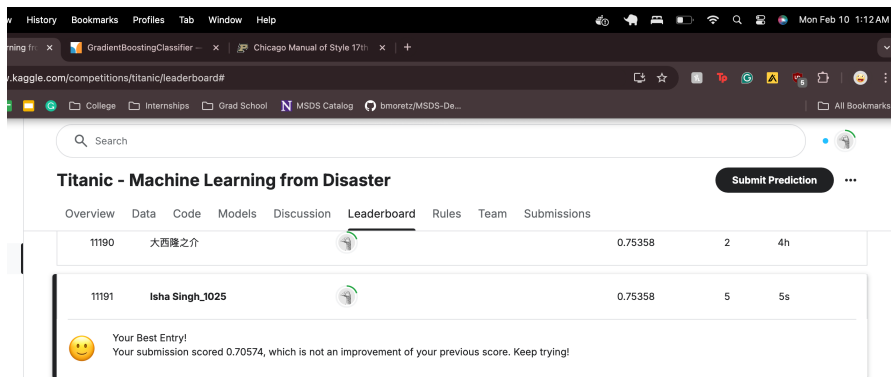
Overview Data Code Models Discussion Leaderboard Rules Team Submissions

11177	Zerast Ibrahim	0.75358	1	2d
11178	Sean Jordan	0.75358	2	1d
11179	Yuchen Lu	0.75358	1	19h
11180	ZengMingYi	0.75358	2	19h
11181	Panchal Shrivharee	0.75358	1	13h
11182	大西隆之介	0.75358	2	3h
11183	Isha Singh_1025	0.75358	4	10s

Your Best Entry!  
Your most recent submission scored 0.75358, which is an improvement of your previous score of 0.54308. Great job!

Tweet this

### Gradient Boosting

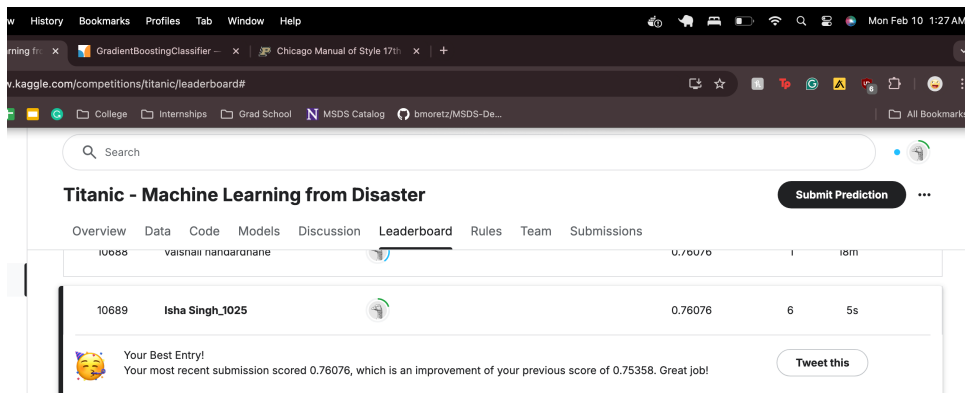


Overview Data Code Models Discussion Leaderboard Rules Team Submissions

11190	大西隆之介	0.75358	2	4h
11191	Isha Singh_1025	0.75358	5	5s

Your Best Entry!  
Your submission scored 0.70574, which is not an improvement of your previous score. Keep trying!

### Extra Trees



Overview Data Code Models Discussion Leaderboard Rules Team Submissions

10689	Isha Singh_1025	0.76076	6	5s
-------	-----------------	---------	---	----

Your Best Entry!  
Your most recent submission scored 0.76076, which is an improvement of your previous score of 0.75358. Great job!

Tweet this

# titanic\_ml module\_05

February 10, 2025

## 1 Titanic - Machine Learning from Disaster

Isha Singh

Professor Irene Tsapara

MSDS 422 Practical Machine Learning

9 February 2025

File Load In

```
[205]: from scipy.stats import shapiro  
       from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[206]: import numpy as np
```

```
[207]: import matplotlib.pyplot as plt
```

```
[208]: import seaborn as sns
```

```
[209]: import pandas as pd
```

```
[210]: from sklearn.ensemble import RandomForestRegressor
```

```
[254]: from scipy.stats import shapiro
```

```
[211]: from sklearn.model_selection import GridSearchCV, KFold, train_test_split
```

```
[270]: from sklearn.metrics import mean_squared_error, r2_score
```

```
[319]: from sklearn.ensemble import GradientBoostingRegressor
```

```
[348]: from sklearn.ensemble import ExtraTreesRegressor
```

SECTION 01: Descriptive Statistics

TRAINING DATA

```
[212]: titanic_train_dataframe = pd.read_csv("/Users/isingh/Desktop/titanic/train.csv")
```

TESTING DATA



```
[213]: titanic_test_dataframe = pd.read_csv("/Users/isingh/Desktop/titanic/test.csv")
```

```
[214]: titanic_test_dataframe.columns
```

```
[214]: Index(['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',  
          'Ticket', 'Fare', 'Cabin', 'Embarked'],  
          dtype='object')
```

```
[215]: titanic_train_dataframe.describe()
```

```
[215]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[216]: missing_values = titanic_train_dataframe.isnull().sum()
```

```
[217]: missing_values
```

```
[217]: PassengerId      0  
Survived           0  
Pclass             0  
Name              0  
Sex               0  
Age              177  
SibSp             0  
Parch             0  
Ticket            0  
Fare              0  
Cabin            687  
Embarked          2  
dtype: int64
```

## Outlier

```
[218]: titanic_train_dataframe["Age"].fillna(  
        titanic_train_dataframe["Age"].median(), inplace=True  
    )
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel\_16658/2219488372.py:1  
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series  
through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work  
because the intermediate object on which we are setting values always behaves as  
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using  
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)  
instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Age'].fillna(titanic_train_dataframe['Age'].median(),  
inplace=True)
```

```
[219]: titanic_train_dataframe["Cabin"].fillna("Unknown", inplace=True)  
titanic_train_dataframe["Embarked"].fillna(  
        titanic_train_dataframe.groupby("Pclass")["Embarked"].transform(  
            lambda x: x.mode().iloc[0]  
        ),  
        inplace=True,  
    )
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel\_16658/2013883459.py:1  
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series  
through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work  
because the intermediate object on which we are setting values always behaves as  
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using  
'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value)  
instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Cabin'].fillna('Unknown', inplace=True)  
/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_16658/2013883459.py:2  
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series  
through chained assignment using an inplace method.  
The behavior will change in pandas 3.0. This inplace method will never work  
because the intermediate object on which we are setting values always behaves as  
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Embarked'].fillna(titanic_train_dataframe.groupby('Pclass')['Embarked'].transform(lambda x: x.mode().iloc[0]), inplace=True)
```

Title

```
[220]: titanic_train_dataframe["Title"] = titanic_train_dataframe["Name"].str.extract(
        " ([A-Za-z]+)\.", expand=False
    )
```

Family Size

```
[221]: titanic_train_dataframe["FamilySize"] = (
        titanic_train_dataframe["SibSp"] + titanic_train_dataframe["Parch"] + 1
    )
```

Whether a family member was alone or not

```
[222]: titanic_train_dataframe["IsAlone"] = (
        titanic_train_dataframe["FamilySize"] == 1
    ).astype(int)
```

Feature Scaling

```
[223]: from sklearn.preprocessing import StandardScaler
```

```
[224]: scaler = StandardScaler()
titanic_train_dataframe[["Age", "Fare"]] = scaler.fit_transform(
    titanic_train_dataframe[["Age", "Fare"]]
)
```

```
[225]: titanic_train_dataframe.duplicated().sum()
```

```
[225]: 0
```

```
[226]: missing_values = titanic_train_dataframe.isnull().sum()
print("Missing Values:")
print(missing_values)
```

Missing Values:

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	0
SibSp	0

```
Parch      0
Ticket     0
Fare       0
Cabin      0
Embarked   0
Title      0
FamilySize 0
IsAlone    0
dtype: int64
```

Outlier IQR

```
[227]: titanic_train_dataframe
```

```
[227]:      PassengerId  Survived  Pclass  \
0             1         0         3
1             2         1         1
2             3         1         3
3             4         1         1
4             5         0         3
..          ...         ...         ...
886          887         0         2
887          888         1         1
888          889         0         3
889          890         1         1
890          891         0         3
```

```
                                Name      Sex      Age  \
0                Braund, Mr. Owen Harris    male -0.565736
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  0.663861
2                Heikkinen, Miss. Laina    female -0.258337
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  0.433312
4                Allen, Mr. William Henry    male  0.433312
..          ...         ...         ...
886                Montvila, Rev. Juozas    male -0.181487
887                Graham, Miss. Margaret Edith    female -0.796286
888  Johnston, Miss. Catherine Helen "Carrie"    female -0.104637
889                Behr, Mr. Karl Howell    male -0.258337
890                Dooley, Mr. Patrick    male  0.202762
```

```
      SibSp  Parch      Ticket     Fare      Cabin Embarked Title  \
0         1      0      A/5 21171 -0.502445   Unknown      S    Mr
1         1      0      PC 17599  0.786845      C85      C   Mrs
2         0      0  STON/O2. 3101282 -0.488854   Unknown      S  Miss
3         1      0      113803  0.420730      C123      S   Mrs
4         0      0      373450 -0.486337   Unknown      S    Mr
..        ...    ...         ...         ...         ...
886        0      0      211536 -0.386671   Unknown      S   Rev
```

887	0	0	112053	-0.044381	B42	S	Miss
888	1	2	W./C. 6607	-0.176263	Unknown	S	Miss
889	0	0	111369	-0.044381	C148	C	Mr
890	0	0	370376	-0.492378	Unknown	Q	Mr

	FamilySize	IsAlone
0	2	0
1	2	0
2	1	1
3	2	0
4	1	1
..	...	...
886	1	1
887	1	1
888	4	0
889	1	1
890	1	1

[891 rows x 15 columns]

```
[228]: numeric_cols = titanic_train_dataframe.select_dtypes(
        include=[np.number]
    ).columns.tolist()
Q1 = titanic_train_dataframe[numeric_cols].quantile(0.25)
Q3 = titanic_train_dataframe[numeric_cols].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filtering out outliers
titanic_train_dataframe1 = titanic_train_dataframe[
    ~(
        (titanic_train_dataframe[numeric_cols] < lower_bound)
        | (titanic_train_dataframe[numeric_cols] > upper_bound)
    ).any(axis=1)
]
```

```
[229]: titanic_train_dataframe1
```

```
[229]: PassengerId  Survived  Pclass  \
0             1         0         3
1             3         1         3
2             4         1         1
3             5         0         3
4             6         0         3
..          ...         ...         ...
```

884	885	0	3
886	887	0	2
887	888	1	1
889	890	1	1
890	891	0	3

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	-0.565736	1	
2	Heikkinen, Miss. Laina	female	-0.258337	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	0.433312	1	
4	Allen, Mr. William Henry	male	0.433312	0	
5	Moran, Mr. James	male	-0.104637	0	
..	...	...	...	...	
884	Sutehall, Mr. Henry Jr	male	-0.335187	0	
886	Montvila, Rev. Juozas	male	-0.181487	0	
887	Graham, Miss. Margaret Edith	female	-0.796286	0	
889	Behr, Mr. Karl Howell	male	-0.258337	0	
890	Dooley, Mr. Patrick	male	0.202762	0	


	Parch	Ticket	Fare	Cabin	Embarked	Title	FamilySize	\
0	0	A/5 21171	-0.502445	Unknown	S	Mr	2	
2	0	STON/O2. 3101282	-0.488854	Unknown	S	Miss	1	
3	0	113803	0.420730	C123	S	Mrs	2	
4	0	373450	-0.486337	Unknown	S	Mr	1	
5	0	330877	-0.478116	Unknown	Q	Mr	1	
..	...	...	...	...	...	...	...	
884	0	SOTON/OQ 392076	-0.506472	Unknown	S	Mr	1	
886	0	211536	-0.386671	Unknown	S	Rev	1	
887	0	112053	-0.044381	B42	S	Miss	1	
889	0	111369	-0.044381	C148	C	Mr	1	
890	0	370376	-0.492378	Unknown	Q	Mr	1	

	IsAlone
0	0
2	1
3	0
4	1
5	1
..	...
884	1
886	1
887	1
889	1
890	1

[577 rows x 15 columns]

Visualizations - Exploratory Data Analysis

```
[231]: plt.figure(figsize=(8, 5))
sns.histplot(titanic_train_dataframe1["Age"], bins=20, kde=True, color="blue")
plt.title("Distribution of Age")
for p in plt.gca().patches:
    plt.gca().annotate(
        f"{int(p.get_height())}",
        (p.get_x() + p.get_width() / 2, p.get_height()),
        ha="center",
        va="bottom",
        fontsize=10,
        color="black",
    )
plt.show()
```




```
[232]: plt.figure(figsize=(8, 5))
sns.barplot(
    x="Pclass",
    y="Survived",
    data=titanic_train_dataframe1,
    palette="coolwarm",
    estimator=np.mean,
)
plt.title("Survival Rate by Passenger Class")
```

```
plt.ylabel("Survival Rate")
plt.show()
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel\_16658/3739617002.py:2  
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Pclass', y='Survived', data=titanic_train_dataframe1,
palette='coolwarm', estimator=np.mean)
```



```
[233]: plt.figure(figsize=(8, 5))
sns.boxplot(x="Pclass", y="Fare", data=titanic_train_dataframe1,
palette="coolwarm")
plt.title("Fare Distribution by Pclass")
plt.show()
```


/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel\_16658/2735142903.py:2  
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same




effect.


```
sns.boxplot(x='Pclass', y='Fare', data=titanic_train_dataframe1,  
palette='coolwarm')
```



```
[234]: numeric_df = titanic_train_dataframe1.select_dtypes(include=[np.number])  
numeric_df = numeric_df.loc[:, (numeric_df.var() > 0)]  
plt.figure(figsize=(12, 8)) # Larger figure size  
sns.heatmap(numeric_df.corr(), annot=True, cmap="coolwarm", fmt=".2f",  
            linewidths=0.5)  
plt.xticks(rotation=45, ha="right")  
plt.yticks(rotation=0)  
plt.title("Feature Correlation Matrix")  
  
plt.show()
```



```
[ ]: numeric_cols = titanic_train_dataframe1.select_dtypes(include=[np.number])
sns.pairplot(numeric_cols, hue="Survived", palette="coolwarm")
plt.suptitle("Scatterplot Matrix", y=1.02)
plt.show()
```



```
[236]: print(titanic_train_dataframe1.isnull().sum())
```

```

PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Cabin          0
Embarked       0

```

```
Title          0
FamilySize     0
IsAlone        0
dtype: int64
```

Train Test Split

```
[238]: from sklearn.model_selection import train_test_split
```

```
X = titanic_train_dataframe1.drop(columns=["Survived"]) # Features
y = titanic_train_dataframe1["Survived"] # Target
```

```
[239]: X = X.drop(columns=["Name", "Ticket", "Cabin", "PassengerId"], errors="ignore")
```

```
[240]: print(X.columns)
X.head()
```

```
Index(['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked', 'Title',
       'FamilySize', 'IsAlone'],
      dtype='object')
```

```
[240]:
```

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title	\
0	3	male	-0.565736	1	0	-0.502445	S	Mr	
2	3	female	-0.258337	0	0	-0.488854	S	Miss	
3	1	female	0.433312	1	0	0.420730	S	Mrs	
4	3	male	0.433312	0	0	-0.486337	S	Mr	
5	3	male	-0.104637	0	0	-0.478116	Q	Mr	

```
FamilySize  IsAlone
0           2        0
2           1        1
3           2        0
4           1        1
5           1        1
```

```
[ ]: X = pd.get_dummies(X, columns=["Sex", "Embarked", "Title"], drop_first=True)
```

```
[242]: X.head()
```

```
[242]:
```

	Pclass	Age	SibSp	Parch	Fare	FamilySize	IsAlone	Sex_male	\
0	3	-0.565736	1	0	-0.502445	2	0	True	
2	3	-0.258337	0	0	-0.488854	1	1	False	
3	1	0.433312	1	0	0.420730	2	0	False	
4	3	0.433312	0	0	-0.486337	1	1	True	
5	3	-0.104637	0	0	-0.478116	1	1	True	

	Embarked_Q	Embarked_S	...	Title_Lady	Title_Major	Title_Master	\
0	False	True	...	False	False	False	
2	False	True	...	False	False	False	

3	False	True	...	False	False	False
4	False	True	...	False	False	False
5	True	False	...	False	False	False

	Title_Miss	Title_Mlle	Title_Mr	Title_Mrs	Title_Ms	Title_Rev	Title_Sir
0	False	False	True	False	False	False	False
2	True	False	False	False	False	False	False
3	False	False	False	True	False	False	False
4	False	False	True	False	False	False	False
5	False	False	True	False	False	False	False

[5 rows x 22 columns]

```
[244]: X = X.astype(int)
X.head()
```

```
[244]:
```

	Pclass	Age	SibSp	Parch	Fare	FamilySize	IsAlone	Sex_male	Embarked_Q	\
0	3	0	1	0	0	2	0	1	0	
2	3	0	0	0	0	1	1	0	0	
3	1	0	1	0	0	2	0	0	0	
4	3	0	0	0	0	1	1	1	0	
5	3	0	0	0	0	1	1	1	1	

	Embarked_S	...	Title_Lady	Title_Major	Title_Master	Title_Miss	\
0	1	...	0	0	0	0	
2	1	...	0	0	0	1	
3	1	...	0	0	0	0	
4	1	...	0	0	0	0	
5	0	...	0	0	0	0	

	Title_Mlle	Title_Mr	Title_Mrs	Title_Ms	Title_Rev	Title_Sir
0	0	1	0	0	0	0
2	0	0	0	0	0	0
3	0	0	1	0	0	0
4	0	1	0	0	0	0
5	0	1	0	0	0	0

[5 rows x 22 columns]

Cross Validation

```
[245]: from sklearn.model_selection import KFold

kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[246]: kf = KFold(n_splits=5, shuffle=True, random_state=42)

for train_index, test_index in kf.split(X, y):
```

```
X_train, X_val = X.iloc[train_index], X.iloc[test_index]
y_train, y_val = y.iloc[train_index], y.iloc[test_index]
```

```
[247]: print(X.dtypes)
print(X.head())
```

```
Pclass      int64
Age          int64
SibSp        int64
Parch        int64
Fare         int64
FamilySize   int64
IsAlone      int64
Sex_male     int64
Embarked_Q   int64
Embarked_S   int64
Title_Dr     int64
Title_Jonkheer int64
Title_Lady   int64
Title_Major  int64
Title_Master int64
Title_Miss   int64
Title_Mlle   int64
Title_Mr     int64
Title_Mrs    int64
Title_Ms     int64
Title_Rev    int64
Title_Sir    int64
dtype: object
```

	Pclass	Age	SibSp	Parch	Fare	FamilySize	IsAlone	Sex_male	Embarked_Q	\
0	3	0	1	0	0	2	0	1	0	
2	3	0	0	0	0	1	1	0	0	
3	1	0	1	0	0	2	0	0	0	
4	3	0	0	0	0	1	1	1	0	
5	3	0	0	0	0	1	1	1	1	

	Embarked_S	...	Title_Lady	Title_Major	Title_Master	Title_Miss	\
0	1	...	0	0	0	0	
2	1	...	0	0	0	1	
3	1	...	0	0	0	0	
4	1	...	0	0	0	0	
5	0	...	0	0	0	0	

	Title_Mlle	Title_Mr	Title_Mrs	Title_Ms	Title_Rev	Title_Sir
0	0	1	0	0	0	0
2	0	0	0	0	0	0
3	0	0	1	0	0	0
4	0	1	0	0	0	0

5            0            1            0            0            0            0

[5 rows x 22 columns]

```
[248]: print("Current columns in X:", X.columns)
```

```
Current columns in X: Index(['Pclass', 'Age', 'SibSp', 'Parch', 'Fare',
'FamilySize', 'IsAlone',
'Sex_male', 'Embarked_Q', 'Embarked_S', 'Title_Dr', 'Title_Jonkheer',
'Title_Lady', 'Title_Major', 'Title_Master', 'Title_Miss', 'Title_Mlle',
'Title_Mr', 'Title_Mrs', 'Title_Ms', 'Title_Rev', 'Title_Sir'],
dtype='object')
```

Model Assumptions

```
[ ]: X_copy = X.copy()
vif_data = pd.DataFrame()
vif_data["Feature"] = X_copy.columns
vif_data["VIF"] = [
    variance_inflation_factor(X_copy.values, i) for i in range(X_copy.shape[1])
]

print("VIF Values for Multicollinearity Check:")
print(vif_data.sort_values(by="VIF", ascending=False))
```

VIF Values for Multicollinearity Check:

	Feature	VIF
5	FamilySize	2271.694262
2	SibSp	376.023709
15	Title_Miss	176.818976
7	Sex_male	135.673565
17	Title_Mr	116.933276
18	Title_Mrs	99.053525
6	IsAlone	12.621587
10	Title_Dr	6.224177
20	Title_Rev	5.063031
12	Title_Lady	3.268537
19	Title_Ms	3.263391
16	Title_Mlle	3.255726
13	Title_Major	3.045242
14	Title_Master	2.038888
21	Title_Sir	2.030270
11	Title_Jonkheer	2.012673
8	Embarked_Q	1.820930
9	Embarked_S	1.747054
0	Pclass	1.240012
1	Age	1.194009
3	Parch	NaN
4	Fare	NaN

```

/Users/isingh/opt/miniconda3/lib/python3.9/site-
packages/statsmodels/regression/linear_model.py:1736: RuntimeWarning: invalid
value encountered in scalar divide
    return 1 - self.ssr/self.centered_tss

```

```

[ ]: numerical_cols = X_copy.select_dtypes(include=[np.number]).columns.tolist()

print("\nShapiro-Wilk Normality Test for LDA:")
for col in numerical_cols:
    try:
        stat, p = shapiro(X_copy[col])
        print(f"- {col}: p-value={p:.4f} {'(Normal)' if p > 0.05 else '(Not_
↪Normal)'}")
    except ValueError:
        print(f"- {col}: Skipped (Shapiro test failed due to large sample_
↪size)")

```

Shapiro-Wilk Normality Test for LDA:

- Pclass: p-value=0.0000 (Not Normal)
- Age: p-value=0.0000 (Not Normal)
- SibSp: p-value=0.0000 (Not Normal)
- Parch: p-value=1.0000 (Normal)
- Fare: p-value=1.0000 (Normal)
- FamilySize: p-value=0.0000 (Not Normal)
- IsAlone: p-value=0.0000 (Not Normal)
- Sex\_male: p-value=0.0000 (Not Normal)
- Embarked\_Q: p-value=0.0000 (Not Normal)
- Embarked\_S: p-value=0.0000 (Not Normal)
- Title\_Dr: p-value=0.0000 (Not Normal)
- Title\_Jonkheer: p-value=0.0000 (Not Normal)
- Title\_Lady: p-value=0.0000 (Not Normal)
- Title\_Major: p-value=0.0000 (Not Normal)
- Title\_Master: p-value=0.0000 (Not Normal)
- Title\_Miss: p-value=0.0000 (Not Normal)
- Title\_Mlle: p-value=0.0000 (Not Normal)
- Title\_Mr: p-value=0.0000 (Not Normal)
- Title\_Mrs: p-value=0.0000 (Not Normal)
- Title\_Ms: p-value=0.0000 (Not Normal)
- Title\_Rev: p-value=0.0000 (Not Normal)
- Title\_Sir: p-value=0.0000 (Not Normal)

```

/Users/isingh/opt/miniconda3/lib/python3.9/site-
packages/scipy/stats/_axis_nan_policy.py:531: UserWarning: scipy.stats.shapiro:
Input data has range zero. The results may not be accurate.
    res = hypotest_fun_out(*samples, **kwargs)

```

Random Forest Regression



```
[279]: param_grid = {
    "n_estimators": [50, 100, 150, 200, 250, 300], # Number of trees
    "max_features": ["sqrt", "log2", 5, 7], # Maximum features per split
    "max_depth": [10, 20, 30, None], # Maximum depth of each tree
    "criterion": ["squared_error", "absolute_error"], # Valid criteria for
    ↪ regression
}
```

Gini and entropy are only valid for classification

Random Forest Regressor is a regression model, which uses different splitting criteria.

“Randomforestclassifier.” scikit. Accessed February 9, 2025. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html>.

```
[280]: rf_regressor = RandomForestRegressor(random_state=42)
```

```
[281]: grid_search_randomforest_regressor = GridSearchCV(
    estimator=rf_regressor,
    param_grid=param_grid,
    cv=5, # 5-fold cross-validation
    scoring="neg_mean_squared_error", # Minimize MSE
)
```

```
[282]: grid_search_randomforest_regressor.fit(X_train, y_train)
```

```
[282]: GridSearchCV(cv=5, estimator=RandomForestRegressor(random_state=42),
    param_grid={'criterion': ['squared_error', 'absolute_error'],
        'max_depth': [10, 20, 30, None],
        'max_features': ['sqrt', 'log2', 5, 7],
        'n_estimators': [50, 100, 150, 200, 250, 300]},
    scoring='neg_mean_squared_error')
```

```
[283]: best_params_random_forest_regressor = grid_search_randomforest_regressor.
    ↪ best_params_
print(
    "Best parameters -- Random Forest Regressor:",
    ↪ best_params_random_forest_regressor
)
```

Best parameters -- Random Forest Regressor: {'criterion': 'squared\_error', 'max\_depth': 10, 'max\_features': 'sqrt', 'n\_estimators': 100}

```
[284]: best_model_with_randomforest = grid_search_randomforest_regressor.
    ↪ best_estimator_
print("Best model with Random Forest Regressor:", best_model_with_randomforest)
```

Best model with Random Forest Regressor: RandomForestRegressor(max\_depth=10, max\_features='sqrt', random\_state=42)

```
[285]: y_pred_random_forest = best_model_with_randomforest.predict(X_val)
```

```
[286]: mse_random_forest_regressor = mean_squared_error(y_val, y_pred_random_forest)
print("The mse for Random Forest Regressor is:", mse_random_forest_regressor)
```

The mse for Random Forest Regressor is: 0.1526500802481171

```
[289]: r_squared_random_forest_regressor = r2_score(y_val, y_pred_random_forest)
print(
    "The r-squared for Random Forest Regressor is:",
    r_squared_random_forest_regressor
)
```

The r-squared for Random Forest Regressor is: 0.23991065087298635

Kaggle Submission for Random Forest Regressor

```
[ ]: X_test = titanic_test_dataframe.drop(
    columns=["Name", "Ticket", "Cabin", "PassengerId"], errors="ignore"
)
```

```
[307]: X_test = pd.get_dummies(X_test, drop_first=True)
```

```
[310]: for col in X_copy.columns:
    if col not in X_test.columns:
        X_test[col] = 0

X_test = X_test[X_copy.columns]
```

```
[311]: X_test.fillna(X_test.median(), inplace=True)
```

```
[312]: print(" X test columns:", X_test.columns)
```

```
X test columns: Index(['Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'FamilySize',
'IsAlone',
    'Sex_male', 'Embarked_Q', 'Embarked_S', 'Title_Dr', 'Title_Jonkheer',
    'Title_Lady', 'Title_Major', 'Title_Master', 'Title_Miss', 'Title_Mlle',
    'Title_Mr', 'Title_Mrs', 'Title_Ms', 'Title_Rev', 'Title_Sir'],
    dtype='object')
```

```
[315]: y_pred_kaggle_randomforest_regression = best_model_with_randomforest.
    predict(X_test)
```

```
[ ]: y_pred_kaggle_binary_conversion_rdf = (
    y_pred_kaggle_randomforest_regression > 0.5
).astype(int)
```

```
[317]: submission_df = pd.DataFrame(
    {
        "PassengerId": titanic_test_dataframe["PassengerId"],
        "Survived": y_pred_kaggle_binary_conversion_rdf,
    }
)
```

```
[318]: submission_file_path = "random_forest_submission.csv"
submission_df.to_csv(submission_file_path, index=False)
```

## Gradient Boosting

“Gradientboostingclassifier.” scikit. Accessed February 10, 2025. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.GradientBoostingClassifier.html>.

```
[372]: param_grid_gradient_boosting = {
    "n_estimators": [50, 150, 300], # number of trees
    "learning_rate": [0.05, 0.1, 0.2], ## Step size per iteration
    "max_depth": [3, 7], # maximum number of levels in each tree
    "subsample": [0.6, 1.0], # Fraction of data used
    "min_samples_split": [2, 10], # Min samples to split
    "min_samples_leaf": [1, 5], # Min samples in leaf
    "max_features": ["sqrt", "log2"], # Features per split
}
```

```
[327]: gradient_boosting_regressor = GradientBoostingRegressor(random_state=42)
```

```
[328]: grid_search_gradient_boosting_regressor = GridSearchCV(
    estimator=gradient_boosting_regressor,
    param_grid=param_grid_gradient_boosting,
    cv=5, # 5-fold cross-validation
    scoring="neg_mean_squared_error",
)
```

```
[329]: grid_search_gradient_boosting_regressor.fit(X_train, y_train)
```

```
[329]: GridSearchCV(cv=5, estimator=GradientBoostingRegressor(random_state=42),
    param_grid={'learning_rate': [0.05, 0.1, 0.2], 'max_depth': [3, 7],
        'max_features': ['sqrt', 'log2'],
        'min_samples_leaf': [1, 5],
        'min_samples_split': [2, 10],
        'n_estimators': [50, 150, 300],
        'subsample': [0.6, 1.0]},
    scoring='neg_mean_squared_error')
```

```
[330]: best_params_gradient_boosting = grid_search_gradient_boosting_regressor.
    ↪ best_params_
print("Best parameters -- Gradient Boosting:", best_params_gradient_boosting)
```

```
Best parameters -- Gradient Boosting: {'learning_rate': 0.1, 'max_depth': 3,
'max_features': 'sqrt', 'min_samples_leaf': 5, 'min_samples_split': 2,
'n_estimators': 50, 'subsample': 1.0}
```

```
[331]: best_model_with_gradient_boosting_regressor = (
        grid_search_gradient_boosting_regressor.best_estimator_
    )
```

```
[334]: y_pred_gradient_boosting = best_model_with_gradient_boosting_regressor.
        ↪predict(X_val)
```

```
[335]: mse_gradient_boosting = mean_squared_error(y_val, y_pred_gradient_boosting)
        print("The mse for Gradient Boosting Regressor is:", mse_gradient_boosting)
```

The mse for Gradient Boosting Regressor is: 0.1410699424712798

```
[337]: r_squared_gradient_boosting = r2_score(y_val, y_pred_gradient_boosting)
        print("The r-squared for Gradient Boosting Regressor is:",
        ↪r_squared_gradient_boosting)
```

The r-squared for Gradient Boosting Regressor is: 0.2975715402173664

Kaggle Submission for Gradient Boosting

```
[340]: X_test_gradient_boosting = titanic_test_dataframe.drop(
        columns=["Name", "Ticket", "Cabin", "PassengerId"], errors="ignore"
    )
```

```
[341]: X_test_gradient_boosting = pd.get_dummies(X_test_gradient_boosting,
        ↪drop_first=True)
```

```
[342]: for col in X_copy.columns:
        if col not in X_test_gradient_boosting.columns:
            X_test_gradient_boosting[col] = 0

        X_test_gradient_boosting = X_test_gradient_boosting[X_copy.columns]
```

```
[343]: X_test_gradient_boosting.fillna(X_test_gradient_boosting.median(), inplace=True)
```

```
[344]: y_pred_kaggle_gradient_boosting = best_model_with_gradient_boosting_regressor.
        ↪predict(
            X_test_gradient_boosting
        )
```

```
[345]: y_pred_kaggle_gradient_boosting = (y_pred_kaggle_gradient_boosting > 0.5).
        ↪astype(int)
```

```
[346]: submission_df_gb = pd.DataFrame(
        {
```

```

        "PassengerId": titanic_test_dataframe["PassengerId"],
        "Survived": y_pred_kaggle_gradient_boosting,
    }
)

```

```

[347]: submission_file_path_gb = "gradient_boosting_submission.csv"
       submission_df_gb.to_csv(submission_file_path_gb, index=False)

```

Extra Trees

“Extratreesclassifier.” scikit. Accessed February 10, 2025. <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.ExtraTreesClassifier.html>.

```

[349]: param_grid_extra_trees = {
        "n_estimators": [50, 150, 300],
        "max_features": ["sqrt", "log2", 0.5],
        "max_depth": [None, 10, 30],
        "min_samples_split": [2, 10],
        "min_samples_leaf": [1, 5],
        "criterion": ["squared_error", "absolute_error"],
    }

```

```

[351]: extra_trees_regressor = ExtraTreesRegressor(random_state=42)

```

```

[ ]: grid_search_extra_trees = GridSearchCV(
    estimator=extra_trees_regressor,
    param_grid=param_grid_extra_trees,
    cv=5, # 5-fold cross-validation
    scoring="neg_mean_squared_error", # Minimize MSE
)

```

```

[355]: grid_search_extra_trees.fit(X_train, y_train)

```

```

[355]: GridSearchCV(cv=5, estimator=ExtraTreesRegressor(random_state=42),
                  param_grid={'criterion': ['squared_error', 'absolute_error'],
                              'max_depth': [None, 10, 30],
                              'max_features': ['sqrt', 'log2', 0.5],
                              'min_samples_leaf': [1, 5],
                              'min_samples_split': [2, 10],
                              'n_estimators': [50, 150, 300]}},
                  scoring='neg_mean_squared_error')

```

```

[356]: best_params_extra_trees = grid_search_extra_trees.best_params_
       print("Best parameters for Extra Trees Regressor:", best_params_extra_trees)

```

```

Best parameters for Extra Trees Regressor: {'criterion': 'squared_error',
'max_depth': 10, 'max_features': 0.5, 'min_samples_leaf': 5,
'min_samples_split': 2, 'n_estimators': 50}

```

```
[358]: best_model_with_extra_trees = grid_search_extra_trees.best_estimator_
```

```
[359]: y_pred_extra_trees = best_model_with_extra_trees.predict(X_val)
```

```
[360]: mse_extra_trees = mean_squared_error(y_val, y_pred_extra_trees)
print("The mse for Extra Trees Regressor is:", mse_extra_trees)
```

The mse for Extra Trees Regressor is: 0.13929092580937663

```
[361]: r2_extra_trees = r2_score(y_val, y_pred_extra_trees)
print("The r-squared for Extra Trees Regressor is:", r2_extra_trees)
```

The r-squared for Extra Trees Regressor is: 0.3064297839499226

```
[362]: X_test_extra_trees = titanic_test_dataframe.drop(
    columns=["Name", "Ticket", "Cabin", "PassengerId"], errors="ignore"
)
```

```
[363]: X_test_extra_trees = pd.get_dummies(X_test_extra_trees, drop_first=True)
```

```
[365]: for col in X_copy.columns:
    if col not in X_test_extra_trees.columns:
        X_test_extra_trees[col] = 0

X_test_extra_trees = X_test_extra_trees[X_copy.columns]
```

```
[366]: X_test_extra_trees.fillna(X_test_extra_trees.median(), inplace=True)
```

```
[368]: y_pred_kaggle_extra_trees = best_model_with_extra_trees.
    ↪predict(X_test_extra_trees)
```

```
[369]: y_pred_kaggle_binary_extra_trees = (y_pred_kaggle_extra_trees > 0.5).astype(int)
```

```
[370]: submission_df_extra_trees = pd.DataFrame(
    {
        "PassengerId": titanic_test_dataframe["PassengerId"],
        "Survived": y_pred_kaggle_binary_extra_trees,
    }
)
```

```
[371]: submission_file_path_extra_trees = "extra_trees_submission.csv"
submission_df_extra_trees.to_csv(submission_file_path_extra_trees, index=False)
```

Evaluation of Models

```
[383]: model_results = pd.DataFrame(
    {
        "Model": ["Random Forest", "Gradient Boosting", "Extra Trees"],
    })
```

```

        "MSE": [mse_random_forest_regressor, mse_gradient_boosting,
↪mse_extra_trees],
        "R_square Score": [
            r_squared_random_forest_regressor,
            r_squared_gradient_boosting,
            r2_extra_trees,
        ],
    }
)

print("Model Performance Summary")
print(model_results.to_string(index=False))

```

```

Model Performance Summary
      Model      MSE  R_square Score
Random Forest 0.152650      0.239911
Gradient Boosting 0.141070      0.297572
    Extra Trees 0.139291      0.306430

```

## MSE COMPARISON

```

[384]: plt.figure(figsize=(10, 7.5))
ax = sns.barplot(x="Model", y="MSE", data=model_results, palette="coolwarm")

for p in ax.patches:
    ax.annotate(
        f"{p.get_height():.4f}",
        (p.get_x() + p.get_width() / 2, p.get_height()),
        ha="center",
        va="bottom",
        fontsize=12,
        color="black",
        fontweight="bold",
    )

plt.ylim(model_results["MSE"].min() - 0.01, model_results["MSE"].max() + 0.02)

# Labels and title
plt.xlabel("Model", fontsize=9)
plt.ylabel("Mean Squared Error (MSE)", fontsize=9)
plt.title("Model Comparison of the Mean Squared Errors", fontsize=12,
↪fontweight="bold")

plt.show()

```

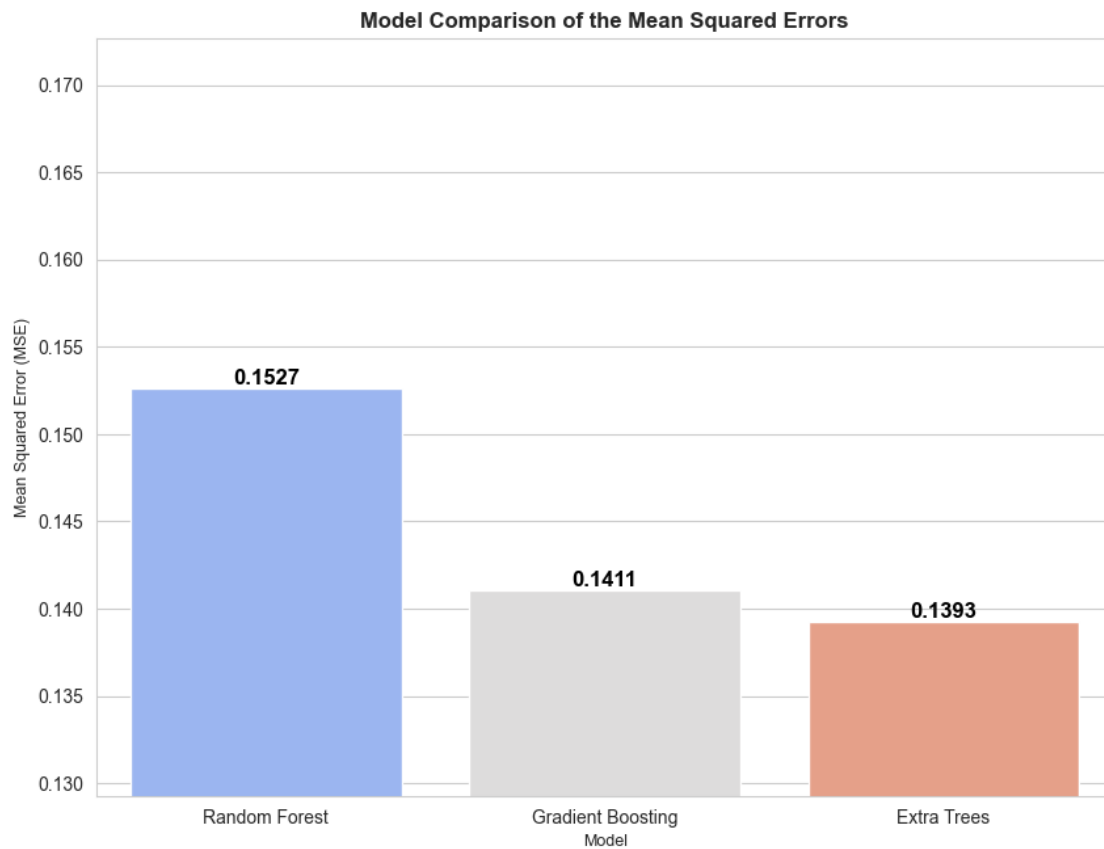
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_16658/632806412.py:2:
FutureWarning:

```

Passing ``palette`` without assigning ``hue`` is deprecated and will be removed in v0.14.0. Assign the ``x`` variable to ``hue`` and set ``legend=False`` for the same effect.

```
ax = sns.barplot(x="Model", y="MSE", data=model_results, palette="coolwarm")
```



R\_Squared

```
[387]: plt.figure(figsize=(10, 7.5))
ax = sns.barplot(x="Model", y="R_square Score", data=model_results,
                palette="coolwarm")

for p in ax.patches:
    ax.annotate(
        f"{p.get_height():.4f}",
        (p.get_x() + p.get_width() / 2, p.get_height()),
        ha="center",
        va="bottom",
        fontsize=12,
        color="black",
        fontweight="bold",
```



```

)

plt.ylim(
    model_results["R_square Score"].min() - 0.05,
    model_results["R_square Score"].max() + 0.05,
)

plt.xlabel("Model", fontsize=9)
plt.ylabel("R_square Score", fontsize=9)
plt.title("Model Comparison of R square Scores", fontsize=12, fontweight="bold")

# Show plot
plt.show()

```


/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel\_16658/679450378.py:2:  
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```

ax = sns.barplot(x="Model", y="R_square Score", data=model_results,
palette="coolwarm")

```



[ ]: