

titanic_ml

February 3, 2025

1 Titanic - Machine Learning from Disaster

Isha Singh

Professor Irene Tsapara

MSDS 422 Practical Machine Learning

2 February 2025

File Load In

```
[236]: from scipy.stats import shapiro  
       from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
[169]: import numpy as np
```

```
[170]: import matplotlib.pyplot as plt
```

```
[171]: import seaborn as sns
```

```
[172]: import pandas as pd
```

SECTION 01: Descriptive Statistics

TRAINING DATA

```
[173]: titanic_train_dataframe=pd.read_csv("/Users/isingh/Desktop/titanic/train.csv")
```

TESTING DATA

```
[174]: titanic_test_dataframe=pd.read_csv("/Users/isingh/Desktop/titanic/test.csv")
```

```
[175]: titanic_test_dataframe.columns
```

```
[175]: Index(['PassengerId', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp', 'Parch',  
        'Ticket', 'Fare', 'Cabin', 'Embarked'],  
        dtype='object')
```

```
[176]: titanic_train_dataframe.describe()
```

```
[176]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[177]: missing_values = titanic_train_dataframe.isnull().sum()
```

```
[178]: missing_values
```

```
[178]: PassengerId      0
Survived              0
Pclass               0
Name                 0
Sex                  0
Age                 177
SibSp                0
Parch                0
Ticket               0
Fare                 0
Cabin               687
Embarked             2
dtype: int64
```

Outlier

```
[179]: titanic_train_dataframe['Age'].fillna(titanic_train_dataframe['Age'].median(),
↳ inplace=True)
```

```
/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/2219488372.py:1
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Age'].fillna(titanic_train_dataframe['Age'].median(),
inplace=True)
```

```
[180]: titanic_train_dataframe['Cabin'].fillna('Unknown', inplace=True)
titanic_train_dataframe['Embarked'].fillna(titanic_train_dataframe.
↳groupby('Pclass')['Embarked'].transform(lambda x: x.mode().iloc[0]),
↳inplace=True)
```

```
/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/2013883459.py:1
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Cabin'].fillna('Unknown', inplace=True)
/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/2013883459.py:2
: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series
through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
titanic_train_dataframe['Embarked'].fillna(titanic_train_dataframe.groupby('Pc
lass')['Embarked'].transform(lambda x: x.mode().iloc[0]), inplace=True)
```

Title

```
[181]: titanic_train_dataframe['Title'] = titanic_train_dataframe['Name'].str.
↳extract(' ([A-Za-z]+)\.', expand=False)
```

Family Size

```
[182]: titanic_train_dataframe['FamilySize'] = titanic_train_dataframe['SibSp'] +  
        ↪titanic_train_dataframe['Parch'] + 1
```

Whether a family member was alone or not

```
[183]: titanic_train_dataframe['IsAlone'] = (titanic_train_dataframe['FamilySize'] ==  
        ↪1).astype(int)
```

Feature Scaling

```
[184]: from sklearn.preprocessing import StandardScaler
```

```
[185]: scaler = StandardScaler()  
        titanic_train_dataframe[['Age', 'Fare']] = scaler.  
        ↪fit_transform(titanic_train_dataframe[['Age', 'Fare']])
```

```
[186]: titanic_train_dataframe.duplicated().sum()
```

```
[186]: np.int64(0)
```

```
[187]: missing_values = titanic_train_dataframe.isnull().sum()  
        print("Missing Values:")  
        print(missing_values)
```

Missing Values:

```
PassengerId    0  
Survived       0  
Pclass         0  
Name           0  
Sex            0  
Age            0  
SibSp          0  
Parch          0  
Ticket         0  
Fare           0  
Cabin          0  
Embarked       0  
Title          0  
FamilySize     0  
IsAlone        0  
dtype: int64
```

Outlier IQR

```
[188]: titanic_train_dataframe
```

```
[188]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	

2	3	1	3
3	4	1	1
4	5	0	3
..
886	887	0	2
887	888	1	1
888	889	0	3
889	890	1	1
890	891	0	3

	Name	Sex	Age	\
0	Braund, Mr. Owen Harris	male	-0.565736	
1	Cummings, Mrs. John Bradley (Florence Briggs Th...	female	0.663861	
2	Heikkinen, Miss. Laina	female	-0.258337	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	0.433312	
4	Allen, Mr. William Henry	male	0.433312	
..	
886	Montvila, Rev. Juozas	male	-0.181487	
887	Graham, Miss. Margaret Edith	female	-0.796286	
888	Johnston, Miss. Catherine Helen "Carrie"	female	-0.104637	
889	Behr, Mr. Karl Howell	male	-0.258337	
890	Dooley, Mr. Patrick	male	0.202762	

	SibSp	Parch	Ticket	Fare	Cabin	Embarked	Title	\
0	1	0	A/5 21171	-0.502445	Unknown	S	Mr	
1	1	0	PC 17599	0.786845	C85	C	Mrs	
2	0	0	STON/O2. 3101282	-0.488854	Unknown	S	Miss	
3	1	0	113803	0.420730	C123	S	Mrs	
4	0	0	373450	-0.486337	Unknown	S	Mr	
..	
886	0	0	211536	-0.386671	Unknown	S	Rev	
887	0	0	112053	-0.044381	B42	S	Miss	
888	1	2	W./C. 6607	-0.176263	Unknown	S	Miss	
889	0	0	111369	-0.044381	C148	C	Mr	
890	0	0	370376	-0.492378	Unknown	Q	Mr	

	FamilySize	IsAlone
0	2	0
1	2	0
2	1	1
3	2	0
4	1	1
..
886	1	1
887	1	1
888	4	0
889	1	1

890 1 1

[891 rows x 15 columns]

```
[189]: numeric_cols = titanic_train_dataframe.select_dtypes(include=[np.number]).
        ↪columns.tolist()
Q1 = titanic_train_dataframe[numeric_cols].quantile(0.25)
Q3 = titanic_train_dataframe[numeric_cols].quantile(0.75)
IQR = Q3 - Q1

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Filtering out outliers
titanic_train_dataframe1 = ↪
        ↪titanic_train_dataframe[~((titanic_train_dataframe[numeric_cols] < ↪
        ↪lower_bound) | (titanic_train_dataframe[numeric_cols] > upper_bound)).
        ↪any(axis=1)]
```

```
[190]: titanic_train_dataframe1
```

```
[190]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
2	3	1	3	
3	4	1	1	
4	5	0	3	
5	6	0	3	
..	
884	885	0	3	
886	887	0	2	
887	888	1	1	
889	890	1	1	
890	891	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	-0.565736	1	
2	Heikkinen, Miss. Laina	female	-0.258337	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	0.433312	1	
4	Allen, Mr. William Henry	male	0.433312	0	
5	Moran, Mr. James	male	-0.104637	0	
..	
884	Sutehall, Mr. Henry Jr	male	-0.335187	0	
886	Montvila, Rev. Juozas	male	-0.181487	0	
887	Graham, Miss. Margaret Edith	female	-0.796286	0	
889	Behr, Mr. Karl Howell	male	-0.258337	0	
890	Dooley, Mr. Patrick	male	0.202762	0	

	Parch		Ticket	Fare	Cabin	Embarked	Title	FamilySize	\
0	0		A/5 21171	-0.502445	Unknown	S	Mr	2	
2	0	STON/O2.	3101282	-0.488854	Unknown	S	Miss	1	
3	0		113803	0.420730	C123	S	Mrs	2	
4	0		373450	-0.486337	Unknown	S	Mr	1	
5	0		330877	-0.478116	Unknown	Q	Mr	1	
..		
884	0	SOTON/OQ	392076	-0.506472	Unknown	S	Mr	1	
886	0		211536	-0.386671	Unknown	S	Rev	1	
887	0		112053	-0.044381	B42	S	Miss	1	
889	0		111369	-0.044381	C148	C	Mr	1	
890	0		370376	-0.492378	Unknown	Q	Mr	1	

	IsAlone
0	0
2	1
3	0
4	1
5	1
..	...
884	1
886	1
887	1
889	1
890	1

[577 rows x 15 columns]


Visualizations - Exploratory Data Analysis

```
[191]: plt.figure(figsize=(8, 5))
sns.stripplot(x='Survived', y='Age', data=titanic_train_dataframe1,
             palette='muted', alpha=0.5, jitter=True)
plt.title('Age Distribution by Survival (Stripplot)')
plt.show()
```


/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/3872080991.py:2
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.stripplot(x='Survived', y='Age', data=titanic_train_dataframe1,
             palette='muted', alpha=0.5, jitter=True)
```



```
[192]: plt.figure(figsize=(8, 5))
sns.histplot(titanic_train_dataframe1['Age'], bins=20, kde=True, color='blue')
plt.title('Distribution of Age')
for p in plt.gca().patches:
    plt.gca().annotate(f'{int(p.get_height())}', (p.get_x() + p.get_width() / 2, p.get_height()),
                      ha='center', va='bottom', fontsize=10, color='black')
plt.show()
```





```
[193]: plt.figure(figsize=(8, 5))
sns.barplot(x='Pclass', y='Survived', data=titanic_train_dataframe1,
           palette='coolwarm', estimator=np.mean)
plt.title('Survival Rate by Passenger Class')
plt.ylabel('Survival Rate')
plt.show()
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/3739617002.py:2
: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Pclass', y='Survived', data=titanic_train_dataframe1,
palette='coolwarm', estimator=np.mean)
```

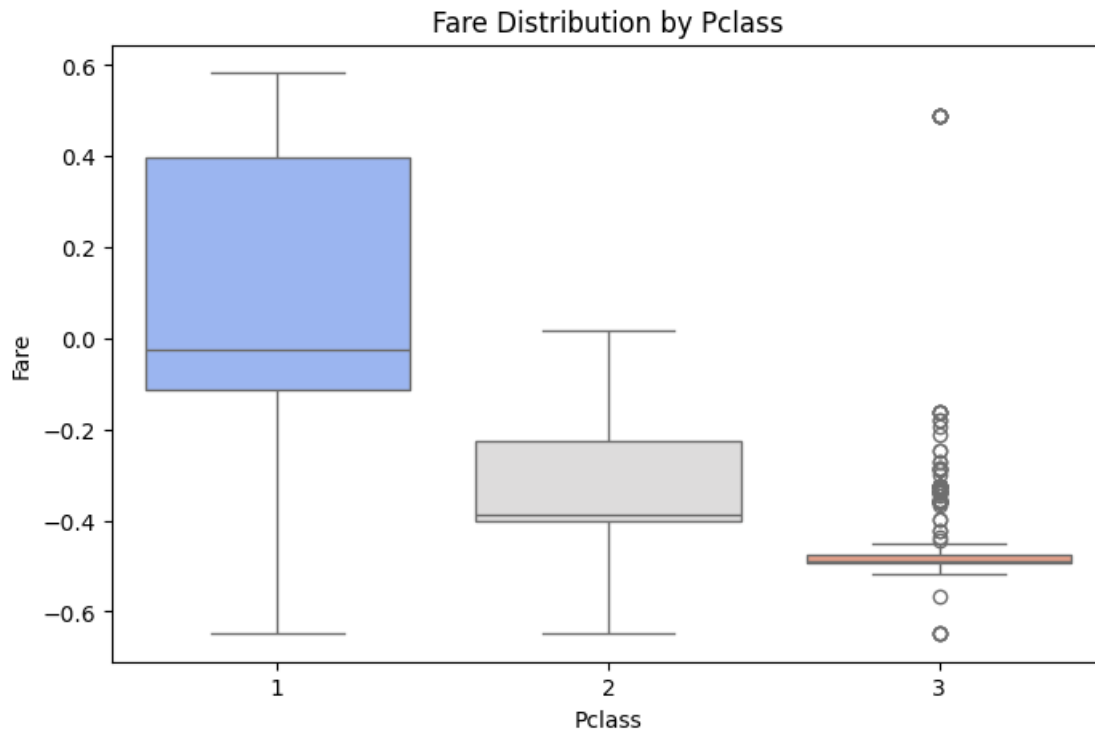


```
[194]: plt.figure(figsize=(8, 5))
sns.boxplot(x='Pclass', y='Fare', data=titanic_train_dataframe1,
           palette='coolwarm')
plt.title('Fare Distribution by Pclass')
plt.show()
```

/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/2735142903.py:2
: FutureWarning:


Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Pclass', y='Fare', data=titanic_train_dataframe1,
palette='coolwarm')
```

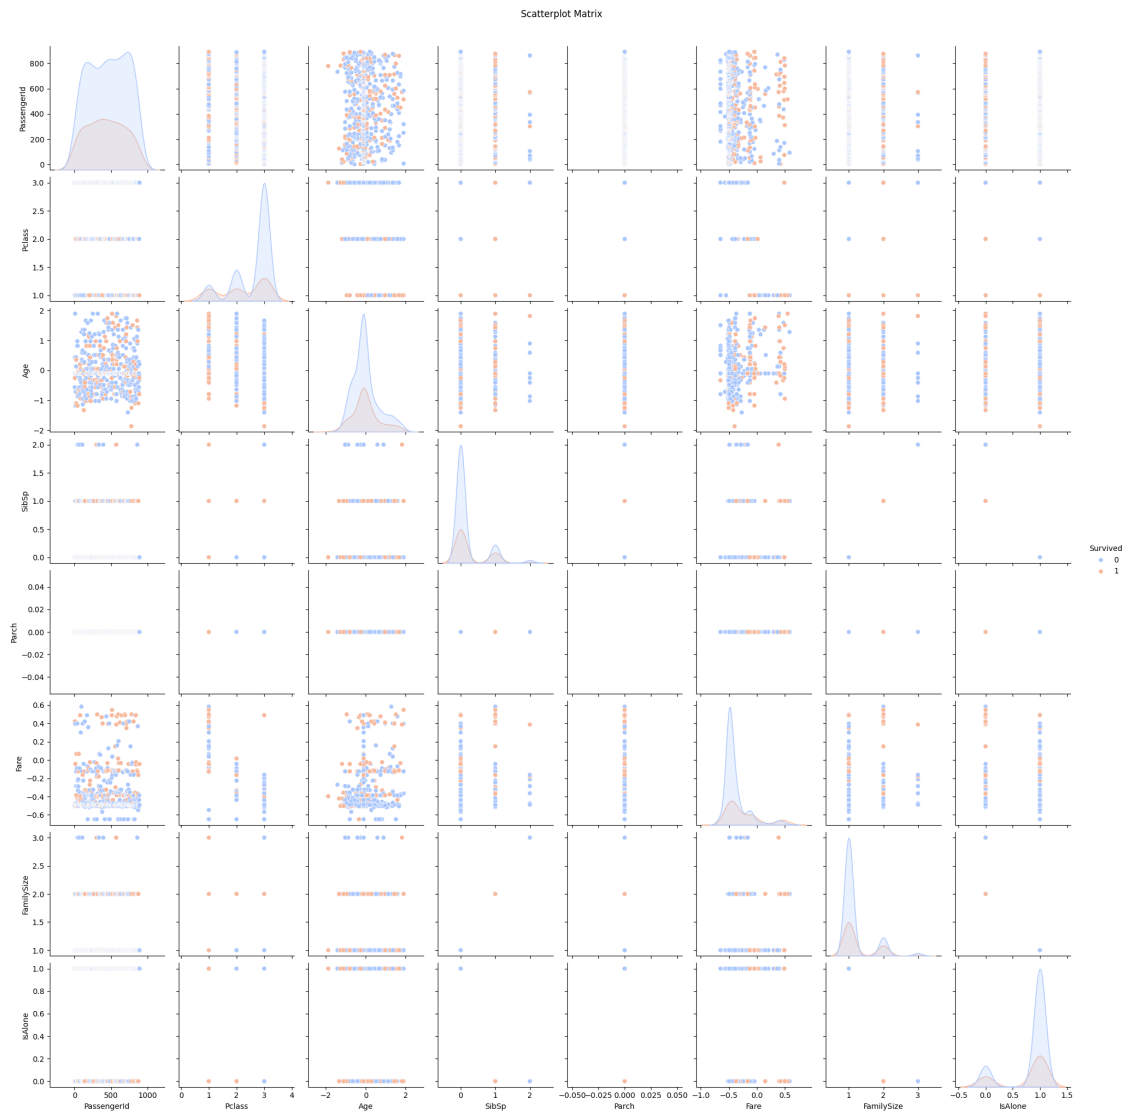


```
[195]: numeric_df = titanic_train_dataframe1.select_dtypes(include=[np.number])
numeric_df = numeric_df.loc[:, (numeric_df.var() > 0)]
plt.figure(figsize=(12, 8)) # Larger figure size
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm', fmt='.2f',
            linewidths=0.5)
plt.xticks(rotation=45, ha='right')
plt.yticks(rotation=0)
plt.title('Feature Correlation Matrix')

plt.show()
```



```
[196]: numeric_cols = titanic_train_dataframe1.select_dtypes(include=[np.number])
sns.pairplot(numeric_cols, hue="Survived", palette="coolwarm")
plt.suptitle('Scatterplot Matrix', y=1.02)
plt.show()
```



```
[197]: print(titanic_train_dataframe1.isnull().sum())
```

```

PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            0
SibSp           0
Parch           0
Ticket         0
Fare           0
Cabin          0
Embarked       0

```

```
Title          0
FamilySize     0
IsAlone        0
dtype: int64
```

Train Test Split

```
[198]: from sklearn.model_selection import train_test_split

X = titanic_train_dataframe1.drop(columns=['Survived']) # Features
y = titanic_train_dataframe1['Survived'] # Target
```

Cross Validation

```
[199]: from sklearn.model_selection import KFold

kf = KFold(n_splits=5, shuffle=True, random_state=42)
```

```
[200]: kf = KFold(n_splits=5, shuffle=True, random_state=42)

for train_index, test_index in kf.split(X, y):
    X_train, X_val = X.iloc[train_index], X.iloc[test_index]
    y_train, y_val = y.iloc[train_index], y.iloc[test_index]
```

```
[201]: X = X.drop(columns=['Name', 'Ticket', 'Cabin', 'PassengerId'], errors='ignore')
```

```
[202]: X = pd.get_dummies(X, columns=['Sex', 'Embarked', 'Title'], drop_first=True)
```

```
[203]: print(X.dtypes)
print(X.head())
```

```
Pclass          int64
Age             float64
SibSp           int64
Parch           int64
Fare            float64
FamilySize      int64
IsAlone         int64
Sex_male        bool
Embarked_Q      bool
Embarked_S      bool
Title_Dr        bool
Title_Jonkheer  bool
Title_Lady      bool
Title_Major     bool
Title_Master    bool
Title_Miss      bool
Title_Mlle      bool
Title_Mr        bool
```

```

Title_Mrs      bool
Title_Ms       bool
Title_Rev      bool
Title_Sir      bool
dtype: object
   Pclass    Age  SibSp  Parch    Fare  FamilySize  IsAlone  Sex_male  \
0        3 -0.565736    1     0 -0.502445         2        0      True
2        3 -0.258337    0     0 -0.488854         1        1     False
3        1  0.433312    1     0  0.420730         2        0     False
4        3  0.433312    0     0 -0.486337         1        1      True
5        3 -0.104637    0     0 -0.478116         1        1      True

   Embarked_Q  Embarked_S  ...  Title_Lady  Title_Major  Title_Master  \
0      False      True  ...      False      False      False
2      False      True  ...      False      False      False
3      False      True  ...      False      False      False
4      False      True  ...      False      False      False
5       True     False  ...      False      False      False

   Title_Miss  Title_Mlle  Title_Mr  Title_Mrs  Title_Ms  Title_Rev  Title_Sir
0      False      False      True   False   False   False   False
2       True      False     False   False   False   False   False
3      False      False     False    True   False   False   False
4      False      False      True   False   False   False   False
5      False      False      True   False   False   False   False

```

[5 rows x 22 columns]

```
[204]: print("Current columns in X:", X.columns)
```

```

Current columns in X: Index(['Pclass', 'Age', 'SibSp', 'Parch', 'Fare',
'FamilySize', 'IsAlone',
'Sex_male', 'Embarked_Q', 'Embarked_S', 'Title_Dr', 'Title_Jonkheer',
'Title_Lady', 'Title_Major', 'Title_Master', 'Title_Miss', 'Title_Mlle',
'Title_Mr', 'Title_Mrs', 'Title_Ms', 'Title_Rev', 'Title_Sir'],
dtype='object')

```

Model Assumptions

KNN

```
[216]: X
```

```

[216]:   Pclass  Age  SibSp  Parch    Fare  FamilySize  IsAlone  Sex_male  \
0        3  22.0    1     0   7.2500         2        0      True
1        1  38.0    1     0  71.2833         2        0     False
2        3  26.0    0     0   7.9250         1        1     False
3        1  35.0    1     0  53.1000         2        0     False
4        3  35.0    0     0   8.0500         1        1      True

```

..
886	2	27.0	0	0	13.0000		1	1	True
887	1	19.0	0	0	30.0000		1	1	False
888	3	28.0	1	2	23.4500		4	0	False
889	1	26.0	0	0	30.0000		1	1	True
890	3	32.0	0	0	7.7500		1	1	True

	Embarked_Q	Embarked_S
0	False	True
1	False	False
2	False	True
3	False	True
4	False	True
..
886	False	True
887	False	True
888	False	True
889	False	False
890	True	False

[891 rows x 10 columns]

```
[221]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import KFold, GridSearchCV, train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score, roc_auc_score,
    roc_curve, precision_recall_curve, auc
)

# Drop unnecessary columns in training data
X = titanic_train_dataframe1.drop(columns=['Survived', 'Name', 'Ticket', 'Cabin', 'PassengerId'], errors='ignore')
y = titanic_train_dataframe1['Survived']

# One-hot encode categorical variables
X = pd.get_dummies(X, drop_first=True)

X.fillna(X.median(), inplace=True)

scaler = StandardScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X), columns=X.columns)
```



```

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
    random_state=42)

kf = KFold(n_splits=5, shuffle=True, random_state=42)

param_grid = {'n_neighbors': np.arange(1, 20)}
knn = KNeighborsClassifier()
grid_search = GridSearchCV(knn, param_grid, cv=kf, scoring='accuracy',
    random_state=42, n_jobs=1) # FIX: n_jobs=1
grid_search.fit(X_train, y_train)
best_knn = grid_search.best_estimator_

accuracies, precisions, recalls, f1_scores, roc_aucs = [], [], [], [], []
tprs, precisions_list = [], []
mean_fpr = np.linspace(0, 1, 100)

for train_index, test_index in kf.split(X_train, y_train):
    X_train_fold, X_val = X_train.iloc[train_index], X_train.iloc[test_index]
    y_train_fold, y_val = y_train.iloc[train_index], y_train.iloc[test_index]

    best_knn.fit(X_train_fold, y_train_fold)
    y_pred = best_knn.predict(X_val)
    y_probs = best_knn.predict_proba(X_val)[:, 1]

    accuracies.append(accuracy_score(y_val, y_pred))
    precisions.append(precision_score(y_val, y_pred, zero_division=1))
    recalls.append(recall_score(y_val, y_pred, zero_division=1))
    f1_scores.append(f1_score(y_val, y_pred, zero_division=1))
    roc_aucs.append(roc_auc_score(y_val, y_probs))

    fpr, tpr, _ = roc_curve(y_val, y_probs)
    tprs.append(np.interp(mean_fpr, fpr, tpr))

    precision, recall, _ = precision_recall_curve(y_val, y_probs)
    precisions_list.append(np.interp(mean_fpr, recall[::-1], precision[::-1]))

# Compute mean ROC and PR AUC
mean_tpr = np.mean(tprs, axis=0)
mean_auc = auc(mean_fpr, mean_tpr)
mean_precision = np.mean(precisions_list, axis=0)
mean_pr_auc = auc(mean_fpr, mean_precision)

plt.figure(figsize=(7, 5))
sns.countplot(x='Survived', data=titanic_train_dataframe1, palette='coolwarm')
plt.title('Survival Count')
for p in plt.gca().patches:

```

```

plt.gca().annotate(f'{{int(p.get_height())}}', (p.get_x() + p.get_width() / 2, p.get_height()),
                  ha='center', va='bottom', fontsize=10, color='black')
plt.show()

# ROC Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_tpr, color='b', label=f"Mean ROC (AUC = {{mean_auc:.2f}})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray", label="Random Classifier")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for KNN Model")
plt.legend()
plt.show()

# Precision-Recall Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_precision, color='r', label=f"Mean PR Curve (AUC = {{mean_pr_auc:.2f}})")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve for KNN Model")
plt.legend()
plt.show()

# Print Model Performance Metrics
print(f"Best Parameters: {{grid_search.best_params_}}")
print(f"Mean Accuracy: {{np.mean(accuracies):.4f}}")
print(f"Mean Precision: {{np.mean(precisions):.4f}}")
print(f"Mean Recall: {{np.mean(recalls):.4f}}")
print(f"Mean F1 Score: {{np.mean(f1_scores):.4f}}")
print(f"Mean ROC AUC: {{np.mean(roc_aucs):.4f}}")

titanic_test_dataframe = pd.read_csv("/Users/isingh/Desktop/titanic/test.csv")

# Store PassengerId separately for submission
passenger_ids = titanic_test_dataframe["PassengerId"]

X_kaggle = titanic_test_dataframe.drop(columns=['Name', 'Ticket', 'Cabin'], errors='ignore')

# One-hot encode categorical variables in test data
X_kaggle = pd.get_dummies(X_kaggle, drop_first=True)

missing_cols = set(X_train.columns) - set(X_kaggle.columns)

```

```

for col in missing_cols:
    X_kaggle[col] = 0

X_kaggle = X_kaggle[X_train.columns]

X_kaggle.fillna(X_kaggle.median(), inplace=True)

X_kaggle_scaled = pd.DataFrame(scaler.transform(X_kaggle), columns=X_kaggle.
    ↪columns)

y_kaggle_preds = best_knn.predict(X_kaggle_scaled)

submission_df = pd.DataFrame({
    "PassengerId": passenger_ids,
    "Survived": y_kaggle_preds
})


submission_df.to_csv("knn_submission.csv", index=False)
print("Submission file 'knn_submission.csv' created. Upload to Kaggle.")


```


/var/folders/qx/htthbr0s1bx5ncc2f9j6j1qc0000gn/T/ipykernel_10058/2957311614.py:6
6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Survived', data=titanic_train_dataframe1, palette='coolwarm')
```







Best Parameters: {'n_neighbors': np.int64(5)}
Mean Accuracy: 0.8069
Mean Precision: 0.7226
Mean Recall: 0.6498
Mean F1 Score: 0.6698
Mean ROC AUC: 0.8266
Submission file 'knn_submission.csv' created. Upload to Kaggle.

Logistic Regression

```
[239]: # Multicollinearity Check (VIF)
vif_data = pd.DataFrame()
vif_data["Feature"] = X_scaled.columns
vif_data["VIF"] = [variance_inflation_factor(X_scaled.values, i) for i in
                    range(X_scaled.shape[1])]

print("VIF Values for Multicollinearity Check:")
print(vif_data.sort_values(by="VIF", ascending=False))
```

VIF Values for Multicollinearity Check:

Feature	VIF
---------	-----

```

2          SibSp          inf
5      FamilySize          inf
15      Title_Miss  177.513865
7          Sex_male  135.483680
17      Title_Mr   117.068989
18      Title_Mrs   99.300740
6          IsAlone   12.850046
10      Title_Dr    6.223599
20      Title_Rev    5.052417
16      Title_Mlle   3.289702
19      Title_Ms     3.271164
12      Title_Lady   3.266631
13      Title_Major   3.023922
4          Fare     2.303734
0          Pclass    2.224417
14      Title_Master  2.042050
11  Title_Jonkheer    2.031291
21      Title_Sir    2.024616
8          Embarked_Q  1.823777
9          Embarked_S  1.729554
1          Age       1.261227
3          Parch          NaN

```

```

/Users/isingh/opt/miniconda3/lib/python3.9/site-
packages/statsmodels/stats/outliers_influence.py:195: RuntimeWarning: divide by
zero encountered in scalar divide

```

```

    vif = 1. / (1. - r_squared_i)

```

```

/Users/isingh/opt/miniconda3/lib/python3.9/site-
packages/statsmodels/regression/linear_model.py:1738: RuntimeWarning: invalid
value encountered in scalar divide

```

```

    return 1 - self.ssr/self.uncentered_tss

```

```

[224]: kf = KFold(n_splits=5, shuffle=True, random_state=42)

```

```

# Logistic Regression Model

```

```

logreg = LogisticRegression(max_iter=500)

```

```

grid_search_logreg = GridSearchCV(logreg, {}, cv=kf, scoring='accuracy',
    ↪n_jobs=1) # No hyperparams to tune

```

```

grid_search_logreg.fit(X_train, y_train)

```

```

best_logreg = grid_search_logreg.best_estimator_

```

```

# Model Evaluation

```

```

accuracies, precisions, recalls, f1_scores, roc_aucs = [], [], [], [], []

```

```

tprs, precisions_list = [], []

```

```

mean_fpr = np.linspace(0, 1, 100)

```

```

for train_index, test_index in kf.split(X_train, y_train):

```

```

    X_train_fold, X_val = X_train.iloc[train_index], X_train.iloc[test_index]

```

```

y_train_fold, y_val = y_train.iloc[train_index], y_train.iloc[test_index]

best_logreg.fit(X_train_fold, y_train_fold)
y_pred = best_logreg.predict(X_val)
y_probs = best_logreg.predict_proba(X_val)[:, 1]

accuracies.append(accuracy_score(y_val, y_pred))
precisions.append(precision_score(y_val, y_pred, zero_division=1))
recalls.append(recall_score(y_val, y_pred, zero_division=1))
f1_scores.append(f1_score(y_val, y_pred, zero_division=1))
roc_aucs.append(roc_auc_score(y_val, y_probs))

fpr, tpr, _ = roc_curve(y_val, y_probs)
tprs.append(np.interp(mean_fpr, fpr, tpr))

precision, recall, _ = precision_recall_curve(y_val, y_probs)
precisions_list.append(np.interp(mean_fpr, recall[::-1], precision[::-1]))

# Compute mean ROC and PR AUC
mean_tpr = np.mean(tprs, axis=0)
mean_auc = auc(mean_fpr, mean_tpr)
mean_precision = np.mean(precisions_list, axis=0)
mean_pr_auc = auc(mean_fpr, mean_precision)

# ROC Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_tpr, color='b', label=f"Mean ROC (AUC = {mean_auc:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray", label="Random Classifier")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for Logistic Regression Model")
plt.legend()
plt.show()

# Precision-Recall Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_precision, color='r', label=f"Mean PR Curve (AUC = {mean_pr_auc:.2f})")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve for Logistic Regression Model")
plt.legend()
plt.show()

```



```

# Print Model Performance Metrics
print(f"Mean Accuracy: {np.mean(accuracies):.4f}")
print(f"Mean Precision: {np.mean(precisions):.4f}")
print(f"Mean Recall: {np.mean(recalls):.4f}")
print(f"Mean F1 Score: {np.mean(f1_scores):.4f}")
print(f"Mean ROC AUC: {np.mean(roc_aucs):.4f}")

titanic_test_dataframe = pd.read_csv("/Users/isingh/Desktop/titanic/test.csv")

# Store PassengerId separately for submission
passenger_ids = titanic_test_dataframe["PassengerId"]

X_kaggle = titanic_test_dataframe.drop(columns=['Name', 'Ticket', 'Cabin'],
    ↪errors='ignore')

X_kaggle = pd.get_dummies(X_kaggle, drop_first=True)

missing_cols = set(X_train.columns) - set(X_kaggle.columns)
for col in missing_cols:
    X_kaggle[col] = 0

X_kaggle = X_kaggle[X_train.columns]

X_kaggle.fillna(X_kaggle.median(), inplace=True)


X_kaggle_scaled = pd.DataFrame(scaler.transform(X_kaggle), columns=X_kaggle.
    ↪columns)


y_kaggle_preds_logreg = best_logreg.predict(X_kaggle_scaled)

submission_df_logreg = pd.DataFrame({
    "PassengerId": passenger_ids,
    "Survived": y_kaggle_preds_logreg
})

submission_df_logreg.to_csv("logreg_submission.csv", index=False)
print("Submission file 'logreg_submission.csv' ")

```





Mean Accuracy: 0.7874
 Mean Precision: 0.6917
 Mean Recall: 0.6019
 Mean F1 Score: 0.6305
 Mean ROC AUC: 0.8102
 Submission file 'logreg_submission.csv'

LDA (Linear Discriminant Analysis)

```
[241]: print("\nShapiro-Wilk Normality Test for LDA:")
for col in numerical_cols:
    stat, p = shapiro(X[col])
    print(f"- {col}: p-value={p:.4f} {'(Normal)' if p > 0.05 else '(Not_
    Normal)'}")
```

Shapiro-Wilk Normality Test for LDA:

- Pclass: p-value=0.0000 (Not Normal)
- Age: p-value=0.0000 (Not Normal)
- SibSp: p-value=0.0000 (Not Normal)
- Parch: p-value=1.0000 (Normal)

- Fare: p-value=0.0000 (Not Normal)
- FamilySize: p-value=0.0000 (Not Normal)
- IsAlone: p-value=0.0000 (Not Normal)

/Users/isingh/opt/miniconda3/lib/python3.9/site-packages/scipy/stats/_axis_nan_policy.py:531: UserWarning: scipy.stats.shapiro: Input data has range zero. The results may not be accurate.

```
res = hypotest_fun_out(*samples, **kws)
```

```
[227]: kf = KFold(n_splits=5, shuffle=True, random_state=42)

lda = LinearDiscriminantAnalysis()
grid_search_lda = GridSearchCV(lda, {}, cv=kf, scoring='accuracy', n_jobs=1) #_
↳ No hyperparameters for LDA
grid_search_lda.fit(X_train, y_train)
best_lda = grid_search_lda.best_estimator_

accuracies, precisions, recalls, f1_scores, roc_aucs = [], [], [], [], []
tprs, precisions_list = [], []
mean_fpr = np.linspace(0, 1, 100)

for train_index, test_index in kf.split(X_train, y_train):
    X_train_fold, X_val = X_train.iloc[train_index], X_train.iloc[test_index]
    y_train_fold, y_val = y_train.iloc[train_index], y_train.iloc[test_index]

    best_lda.fit(X_train_fold, y_train_fold)
    y_pred = best_lda.predict(X_val)
    y_probs = best_lda.predict_proba(X_val)[:, 1]

    accuracies.append(accuracy_score(y_val, y_pred))
    precisions.append(precision_score(y_val, y_pred, zero_division=1))
    recalls.append(recall_score(y_val, y_pred, zero_division=1))
    f1_scores.append(f1_score(y_val, y_pred, zero_division=1))
    roc_aucs.append(roc_auc_score(y_val, y_probs))

    fpr, tpr, _ = roc_curve(y_val, y_probs)
    tprs.append(np.interp(mean_fpr, fpr, tpr))

    precision, recall, _ = precision_recall_curve(y_val, y_probs)
    precisions_list.append(np.interp(mean_fpr, recall[::-1], precision[::-1]))

# Compute mean ROC and PR AUC
mean_tpr = np.mean(tprs, axis=0)
mean_auc = auc(mean_fpr, mean_tpr)
mean_precision = np.mean(precisions_list, axis=0)
mean_pr_auc = auc(mean_fpr, mean_precision)
```

```

# ROC Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_tpr, color='b', label=f"Mean ROC (AUC = {mean_auc:.2f})")
plt.plot([0, 1], [0, 1], linestyle="--", color="gray", label="Random Classifier")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve for LDA Model")
plt.legend()
plt.show()

# Precision-Recall Curve
plt.figure(figsize=(8, 6))
plt.plot(mean_fpr, mean_precision, color='r', label=f"Mean PR Curve (AUC = {mean_pr_auc:.2f})")
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve for LDA Model")
plt.legend()
plt.show()

# Print Model Performance Metrics
print(f"Mean Accuracy: {np.mean(accuracies):.4f}")
print(f"Mean Precision: {np.mean(precisions):.4f}")
print(f"Mean Recall: {np.mean(recalls):.4f}")
print(f"Mean F1 Score: {np.mean(f1_scores):.4f}")
print(f"Mean ROC AUC: {np.mean(roc_aucs):.4f}")

titanic_test_dataframe = pd.read_csv("/Users/isingh/Desktop/titanic/test.csv")

passenger_ids = titanic_test_dataframe["PassengerId"]

X_kaggle = titanic_test_dataframe.drop(columns=['Name', 'Ticket', 'Cabin'], errors='ignore')

X_kaggle = pd.get_dummies(X_kaggle, drop_first=True)

missing_cols = set(X_train.columns) - set(X_kaggle.columns)
for col in missing_cols:
    X_kaggle[col] = 0

X_kaggle = X_kaggle[X_train.columns]

X_kaggle.fillna(X_kaggle.median(), inplace=True)

```

```

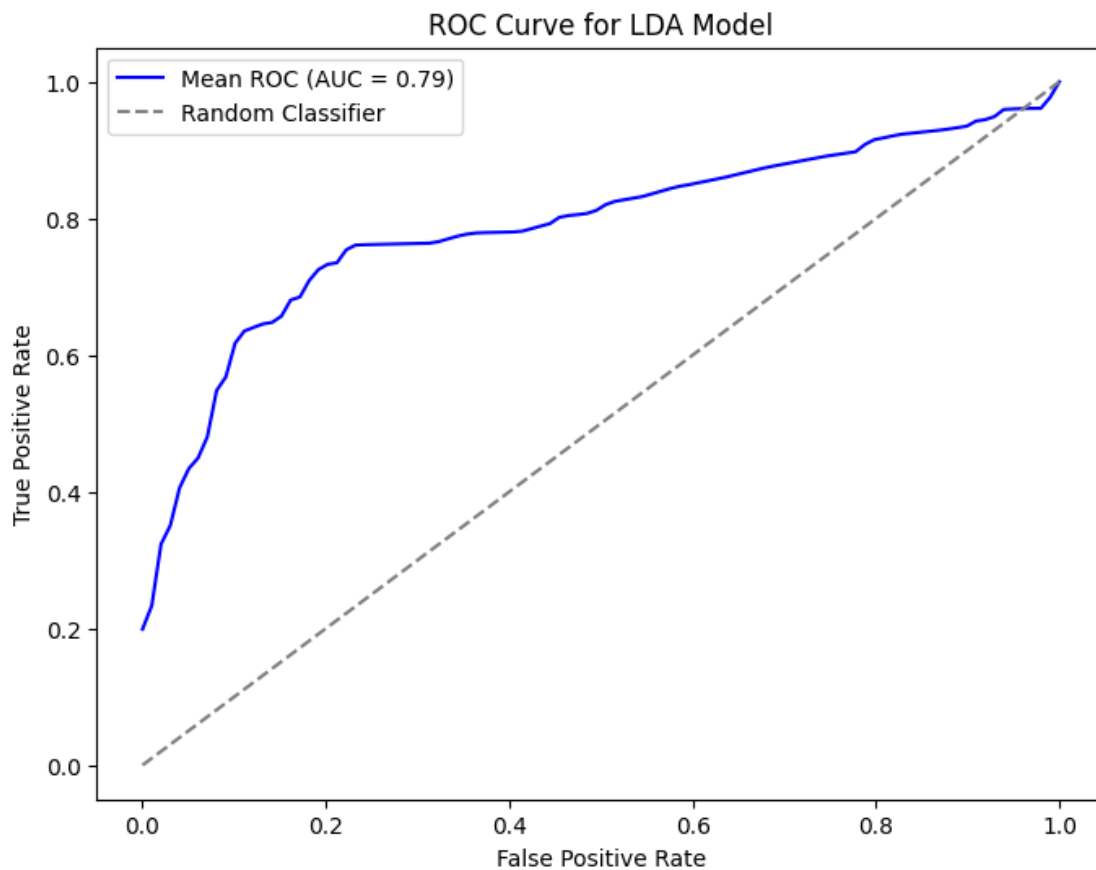
# Standardize test data
X_kaggle_scaled = pd.DataFrame(scaler.transform(X_kaggle), columns=X_kaggle.
    ↪columns)

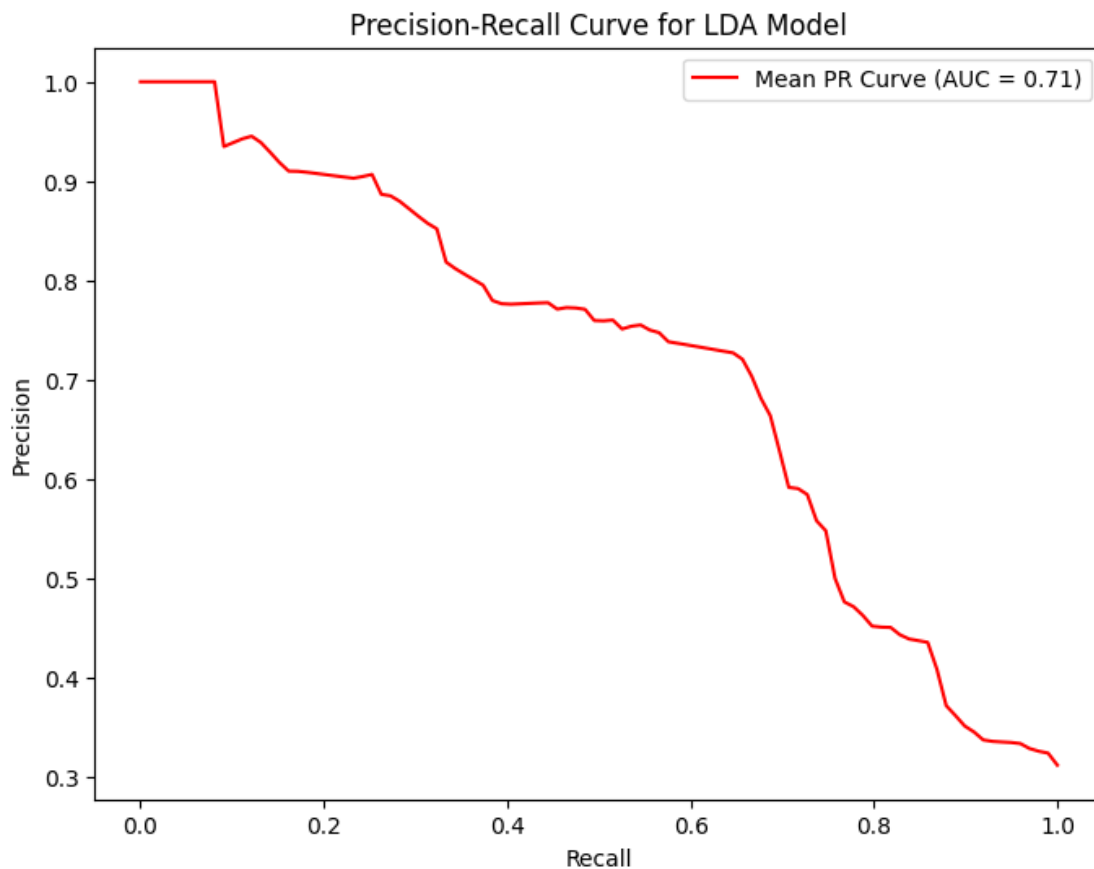
y_kaggle_preds_lda = best_lda.predict(X_kaggle_scaled)

submission_df_lda = pd.DataFrame({
    "PassengerId": passenger_ids,
    "Survived": y_kaggle_preds_lda
})

submission_df_lda.to_csv("lda_submission.csv", index=False)
print("Submission file 'lda_submission.csv' created. Upload to Kaggle.")

```





Mean Accuracy: 0.8005

Mean Precision: 0.7178

Mean Recall: 0.6227

Mean F1 Score: 0.6549

Mean ROC AUC: 0.7888

Submission file 'lda_submission.csv' created. Upload to Kaggle.

[]: